

Joint Granular Model for Load, Solar and Wind Power Scenario Generation

René Carmona, *Member IEEE*, & Xinshuo Yang

Abstract—The main thrust of the paper is the development of a joint stochastic model for electricity demand, and wind and solar power production in a given region. The model hinges on special statistical data analysis techniques including the estimation of heavy tail distributions, graphical LASSO fitting procedures, and conditional Monte Carlo simulations. Assuming the availability of point forecasts, we model the deviations from these forecasts instead of modeling the actual quantities of interest. The resolution of the model is determined by the resolution of the forecast data. For the sake of illustration, we implement our model and the corresponding simulation algorithms on data made available by NREL for the Texas region with hourly time resolution, load data at the zone level, and wind and solar power production at the generation asset level. Our numerical simulations confirm that the dependencies identified through the fitting algorithm are consistent with the relative locations of the production assets and the geographical load zones over which the data were collected.

Index Terms—Load, Solar Power, Wind power, Monte Carlo simulations, Graphical LASSO, Conditional Simulation.

I. INTRODUCTION

MODERN operation of the electric grid has become highly technical, but despite clear evidence of the stochasticity of electricity demand, the major unit commitment and economic dispatch computer algorithms remained intrinsically deterministic. This is even more problematic given the increased penetration of renewable production sources like solar and wind whose uncertainty and volatility have been universally documented. For this reason, it became commonplace to have engineers rely on probabilistic forecasts of wind, solar and even loads, in lieu of deterministic point forecasts. Even though these new classes of forecasts provide a form of quantification of the fluctuations of the variables, and hence of the reliability of the point forecasts, they fail to capture temporal and spatial dependencies between the random quantities they try to predict. Still they became the sources of a very large number of publications. For the methodologies behind these probabilistic forecasts we refer to the following recent surveys: [1] for load, [2] for wind power, and [3] for solar power generation. While the application domain is the same, our approach differs significantly from those papers. Indeed, we do not aim at producing forecasts. Instead, we assume the availability of historical data for both actuals and their forecasts which we use to model the deviations between the two. This choice is motivated by the fact that many users, especially ISOs, commissioned their own forecasts which they do not want to part from, whether they purchased them, or they

developed them in house. Our stance on the matter is more towards *risk management* rather than *prediction*. Also, we demonstrate in the paper how the quality of our scenarios may depend upon the quality of the forecasts we use, a point which is clearly made in the study [4] quantifying the economic value of forecast improvement. Modeling the random behavior of the future values of a variable, say wind, is commonly done on the basis of existing forecasts obtained exogenously. See for example [5]. This trend has become a standard in the industry, and we shall follow this time honored convention in the present paper.

A. Background.

Statistical characterization of wind power output using time series analysis has been the object of numerous publications for over 15 years. See for example [5]–[10]. The existing literature contains many attempts to include spatial and temporal dependencies in models designed to produce forecasts and/or be implemented for Monte Carlo scenario generation. For example, [10] uses multivariate ARMA models to capture temporal correlations, the variance–covariance matrix characterizing the spatial correlations of wind speeds. In the same spirit, [6] proposes ARMAX time series models combined with LASSO estimation, and applies the model to a set of 50 wind farms in Brazil. In the same spirit, [7] proposes to use ARIMA-GARCH models for the time evolution of wind power plant production, and to build dependencies between them by using hierarchical copulas based on two-dimensional copulas like t , Clayton and Gumbel copulas. The authors report on a case study involving 8 wind farms. Similar ideas based on univariate VAR-mGARCH models are used in [8] to model a small number of dependent univariate time series. In another case study of the Brazilian electricity market, and with a slightly different focus on economic costs of the joint management of a wind farm and a hydro generation asset, [9] uses Monte Carlo scenarios generated from a joint model *wind-hydro* to quantify the benefits of adding a renewable source to a hydro facility. From a slightly different point of view, [11] uses a special form of generalized dynamic factor models based on power spectral density estimates and vector auto-regressive models to jointly model load and wind power generation. The scale-up of these methods to accommodate hundreds of wind farms and solar stations is the major shortcoming of these methods and one of the reasons we propose an alternative approach in this paper.

Among the many attempts at extending the realm of time series analysis to include possible dependencies between the

This paper was produced by the IEEE Publication Technology Group. They are in Piscataway, NJ.

Manuscript received April 19, 2021; revised August 16, 2021.

components of a high-dimensional time series, hierarchical forecasting received a lot of attention. See for example the book [12] and the references therein. Two serious impediments prevented us from implemented such methods: 1) the high dimensionality of the *geographical component* when compared to the short length of the historical data, and 2) the fact that our goal goes beyond the production of forecasts, since forecasts are part of the data we work from.

Another important line of research devoted to the generation of spatial-temporal wind power scenarios is proposed in the seminal papers [13]–[17]. While limited to the modeling of a homogeneous ensemble of wind power stations, [15] may be the closest one, at least in spirit, to the approach proposed in this paper. In [15], working in a Gaussian copula framework, the authors analyzed the wind forecast errors and proposed intricate parameterizations of the precision matrices to capture the underlying spatial-temporal correlation structure. Unfortunately, this approach requires large numbers of samples to accurately capture the interdependence structure between the forecast errors.

Finally, with a clear goal of providing inputs for unit commitment and economic dispatch platforms, a recent work in [18] explores a non-parametric method to employ epispine basis functions to approximates the stochastic process for renewable energy and controls the degree to which extreme errors are captured. This approach has been quite successful when assessed using a number of qualitative scenario evaluation metrics. For a comprehensive discussion of the evaluation of the quality of scenarios, we refer to the fundamental paper [19] and the reference therein.

B. Contributions of the paper.

Most of the highly sophisticated optimization programs providing unit commitment and economic dispatch for the operation of modern electricity grids remain intrinsically deterministic despite the widely acknowledged need for the introduction of stochastic factors. To remedy this shortcoming, we propose an original joint model for electricity demand, and wind and solar power generation capturing spatial and temporal dependencies between the renewable generation assets and load zones. The originality of our contribution to the literature is the inclusion of spatial and temporal dependencies at the production asset granularity level, as well as the implementation of Monte Carlo simulation algorithms producing scenarios which can be used for applications ranging from risk assessment to future generation asset siting.

Before delving into the gory details of our modeling approach and the inner working of our simulation engines, we reiterate that we follow the time honored approach of modeling the deviations from forecasts. The down side is that the historical data needed to fit our model needs to include actual measurements of load, wind and solar power as well as their respective forecasts. Our models and their implementations as simulation engines are agnostic to the source of the forecasts. However, the quality of the latter does impact the goodness of fit of our models as well as the predictive power of the Monte Carlo simulator.

We now describe in broad brushstrokes the several prongs of our approach, in no small part to emphasize the originality of our contribution by contrast to the existing literature. In some respect, our approach is reminiscent of the Gaussian copula approach as advocated for example in [13] and [16]. However, it differs in significant ways. First, we recognize that the marginal distributions of the hourly zonal load deviations have heavy tails. Therefore, we fit Generalized Pareto Distributions (GPDs) to the marginals whenever heavy tails are detected. We use them to process the data in order to give them Gaussian marginals. Then, instead of relying on the parametric forms used in [16], we use non-parametric estimation techniques based on the LASSO L^1 regularization to estimate large covariance matrices. We explain and illustrate the pros and cons of this step with simulation results on days for which the forecasts overestimate and underestimate the actuals. Additionally, because we want our Monte Carlo simulations to be strictly non-anticipative, and not *use a crystal ball* when it comes to the generation of scenarios for the future, we use only historical data prior to the day for which we need scenarios. As a result, on some days, one may only have a small number of historical observations when fitting the model. In such cases, the large dimension of the state variables of the model is likely to result in singular empirical covariance matrix estimates. Such a situation is often encountered in modern high dimensional statistics, and in order to overcome this shortcoming, we posit (after checking that this assumption is consistent with the data) that the covariance and precision matrices are separable in the sense that they are tensor products of smaller matrices. This makes it easier to disentangle the contributions of the spatial and temporal components to the correlation structures. This reduced form of the model can then be estimated using the *gemini* algorithm which is tailored to the case of separable covariance structures, and which was introduced in [20] as a variation on the graphical lasso algorithm.

The model introduced in this paper is built from the ground up. We first introduce marginal models separately for the three features: load, wind power and solar power. Historical data for load actuals and forecasts being only available at the zone level, possible correlations between load and wind on one hand, and between load and solar on the other, need to be detected and implemented at the load-zone level, even if eventually, wind and solar power need to be modeled and simulated at the individual asset level.

The marginal model for load at the zone level is inspired by our earlier work [21] on the ERCOT data publicly available from the Texas ISO web site. While the load data used in the present paper is very different from the data used in [21], load data is still at the zone level, and even though the zones are different in their number and coverage, the modeling approach used in the present paper is similar to the one of our previous work. The stochastic models we developed for load, and wind and solar power are the fulcrum for a Monte Carlo simulation engine providing scenarios on demand. Given a date and a set of historical observations (including forecasts) prior to that date, the simulation engine produces as many Monte Carlo scenarios as requested for the next 24 hourly values of load,

wind and solar,

Our simulation engine outputs scenarios in the format predicated by the format of the historical data used to fit the models. It is implemented in a Python package made available by Michal Grzadkowski with extensive README files and examples which can be run from Jupyter Notebooks. It can be found at [22].

The proposed methodology provides a general framework to model the spatial and temporal correlations between load and renewable production forecast deviations. It comprises two major building blocks, model fitting and scenario generation. The model fitting procedure includes:

- 1) Fitting a model to historical data based on marginal distribution estimation and a high dimensional Gaussian copula:
 - a) For each time horizon and location (load zone, wind farm or solar station), fitting a marginal distribution function (e.g., Generalized Pareto Distribution) to the historical data.
 - b) Applying probability integral transforms to transform the data so it becomes uniformly distributed on $[0, 1]$, and then Gaussian $N(0, 1)$.
- 2) For solar, due to its diurnal nature, substitute coefficients in an eigenvector basis computed from a Principal Component Analysis (PCA) procedure to the original hourly data.
- 3) Estimating a separable correlation structure of the Gaussian data using the GEMINI variant of graphical LASSO.

The scenario generation engine involves:

- 1) Generating Monte Carlo samples from the estimated multivariate Gaussian distribution. This involves conditional simulations to reverse the aggregation steps.
- 2) For solar, applying the inverse PCA transform to the samples.
- 3) For each location (load zone, wind farm or solar station) and time horizon, reversing the steps in the Gaussian copula to obtain samples with the appropriate marginal distributions.
- 4) Adding back the forecasts to obtain the desired Monte Carlo scenarios.

C. Organization of the paper.

The paper is organized as follows: in Section II we describe the data used to build an instance of the stochastic model and produce illustrations. The data was assembled by NREL for the Department of Energy program ARPA-E PERFORM, [23]. It covers the Texas region, but we also ran instances of the model for the New York ISO region and MISO. Section III is devoted to the analysis of the aggregate demand for electricity at the zone level. The fitted model is the basis for the Monte Carlo simulation engine whose performance is demonstrated by examples of batches of scenarios. Our asset level model is developed and fitted to the wind power generation historical data in Section IV. In Section V, we introduce a marginal model for the 24 hours of solar power production for each of the solar stations for which we have historical data. Principal

Component Analysis (PCA) is used to circumvent some of the technical difficulties created by solar irradiation idiosyncrasies. The three marginal models are blended into a joint model in Section VI using aggregation procedures and conditional simulation techniques. We conclude with a short recall of the main contributions of the paper in Section VIII.

II. PROBLEM SET-UP AND DESCRIPTION OF THE DATA

For the purpose of illustration, we use the Texas data assembled by NREL for the Department of Energy program ARPA-E PERFORM [23].¹ This synthetic data set contains time series for electricity demand at the zone level, and actuals and forecasts of wind and solar power generation at the asset level. Several time frequencies are provided, but for the sake of definiteness, we shall limit ourselves to hourly data. The geographical resolution is the highest in the sense that data is provided at the site-level.

The **solar** data comprises 226 sites (including 22 existing and 204 proposed solar sites) in the Texas region. The solar power forecasts are provided for 2018 as probabilistic forecasts in the form of 1 to 99 percentiles. These probabilistic forecasts are generated using a Bayesian Model Averaging (BMA) [23] procedure to the ensembles of 51 deterministic forecasts provided by the European Centre for Medium-Range Weather Forecasts (ECMWF). However, since the ECMWF forecasts included in the NREL data set are for 2017 as well as 2018, in order to maximize the size of historical data to which we fit our model, we chose to ignore the NREL probabilistic forecasts and use instead, a plain average of the 51 ECMWF deterministic forecasts as a point forecast. In summary, the day-ahead forecasts we use have a 12-hour-ahead lead time, a 24-hour horizon, an hourly resolution, and a daily update rate. For each day, at 12:00 (issue time), we have a set of hourly forecasts for the next day starting from 00:00 ending at 23:00 (forecast time). For the actual solar power, the data is given with a 5-min resolution for 2017 and 2018, so we computed an hourly aggregate.

The **wind** data comprises 264 sites (including 125 existing and 139 proposed sites) in the Texas region. The wind day-ahead forecasts are provided by NREL in the same format as the solar data – probabilistic forecasts for 2018 and 51 ECMWF forecasts for 2017 and 2018. However, the ECMWF forecasts are highly biased and point forecasts obtained by plain averaging yield much larger deviations. For this reason, we chose to use the median of the probabilistic forecasts as a point forecast. In order to demonstrate the impact of the quality of the forecasts on the performance of the scenario generation, we also use the plain average of the 51 ECMWF forecasts for which the available historical data covers 2017 and 2018 while historical data is limited to 2018 if we use the median of the probabilistic forecasts.

The NREL data set also contains **load** data at the zone level. Actuals and forecasts were collected from the ERCOT official website at hourly resolution for 8 load zones. The forecasts

¹We implemented our model on several synthetic grids, including the NY-ISO region [23] and the IEEE's Reliability Test System (RTS) [24]. Detailed examples can be found at <https://github.com/PrincetonUniversity/PGScen>.

were obtained from ERCOT's "Seven-Day Load Forecast by Weather Zone", see [25], by selecting the data issue and forecast times to match those of the solar day-ahead forecasts.

III. MARGINAL MODEL FOR ELECTRICITY LOADS

A. Strategy

The data covers the period 2017-01-01 to 2018-12-31 for $N_Z = 8$ load zones. For each day d and each load zone $z \in \mathcal{Z} = \{\text{Coast, East, FarWest, North, NorthCentral, South, SouthCentral, West}\}$, we have a vector of $N_h = 24$ load point forecasts, the h -th element representing the point forecast at time horizon h , for $h \in \{0, 1, \dots, N_h - 1\}$. If we let the load point forecast be $\tilde{L}_{z,h}^d$ and the corresponding actual be $\bar{L}_{z,h}^d$, we model the deviations $L_{z,h}^d = \bar{L}_{z,h}^d - \tilde{L}_{z,h}^d$ (often called errors in the literature on the subject).

1) For each load zone z and for each time horizon $h \in \{0, 1, \dots, N_h - 1\}$ the load deviations $L_{z,h}^d$ form a daily time series of length N_d , the number of days prior to the day we generate scenarios for. The first step of our analysis is to analyze the marginal distributions of these $N_Z \times N_h = 8 \times 24 = 192$ stationary univariate time series.

2) For each load zone z and for each time horizon h we produce a Q-Q plot of the empirical quantiles of the load deviations $\{L_{z,h}^d\}_{d=1}^{N_d}$ against the theoretical quantiles of the standard Gaussian distribution $N(0, 1)$. Figure 1 give a sample of 2 of these Q-Q plots. They are typical of what we found throughout. This is clear evidence that the marginal distributions of the load deviations $\{L_{z,h}^d\}_{d=1}^{N_d}$ have *heavy tails*, justifying our next step.




Fig. 1. Normal Q-Q plots for load deviation remainders in the $z = \text{Coast}$ region and time horizons $h = 1$ and 3 hours ahead.

3) For each load zone z and for each time horizon h we fit a Generalized Pareto Distribution (GPD), see for example [26, Chapter 2], to the load deviations $\{L_{z,h}^d\}_{d=1}^{N_d}$. We denote such a distribution as $G_{z,h}$ for the sake of later reference. Throughout this work, in order to manipulate GPD distributions, we use functions of the R library `Rsaofd` [27] accompanying the book [26], which we wrapped into a Python package.

4) For each load zone z and for each time horizon h we transform the load deviation time series $\{L_{z,h}^d\}_{d=1}^{N_d}$ into a uniform time series by the *probability integral transform*. To be specific, let the cumulative distribution function (CDF) of the GPD be $\Phi_{G_{z,h}}$, for every d , we compute $\Phi_{G_{z,h}}(L_{z,h}^d)$.

If $G_{z,h}$ is a good fit, we should expect the marginal distribution of the new time series to be uniform. The histogram on the left pane of Figure 2 is a testament to this claim.




Fig. 2. Transformations of the load deviations in the Coast region 4 hours ahead.

5) For each load zone z and for time horizon h we transform the uniform time series $\{\Phi_{G_{z,h}}(L_{z,h}^d)\}_{d=1}^{N_d}$ into a time series with a standard Gaussian marginal distribution $N(0, 1)$. In order to do that, we use the *inverse probability integral transform*. For each d , we compute $\hat{L}_{z,h}^d = \Phi^{-1}(\Phi_{G_{z,h}}(L_{z,h}^d))$ where Φ^{-1} is the quantile function (inverse CDF) of the standard Gaussian distribution.

6) At this stage of the analysis, we have obtained a data matrix which contains N_d rows and $N_Z \times N_h = 8 \times 24 = 192$ columns with each column following a standard Gaussian distribution. While we do not know if this $N_Z \times N_h$ -dimensional random variables are actually *jointly Gaussian*, we shall behave as if that was the case. In other words, we assume the $N_Z \times N_h$ -dimensional random variables follow a standard Gaussian distribution with mean zero and a covariance matrix Σ , i.e., $\hat{L}^d \sim N(0, \Sigma)$, where $\hat{L}^d$ denotes the d -th row in the data matrix. Given the high dimensionality of the random variables when compared to the small amount of historical data available to estimate the covariance structure of such a Gaussian vector, we expect that any empirical estimate of the covariance matrix will be degenerate and hence not invertible. Therefore, in order to estimate the precision matrix capturing the dependencies between the various components of the vector, we use a regularization procedure aiming at a sparse estimator: we fit a *LASSO Gaussian Graphical Model*. The graphical LASSO estimator of $\Theta = \Sigma^{-1}$ is:

$$\hat{\Theta} = \arg \min_{\Theta \geq 0} (\text{tr}(S\Theta) - \log \det(\Theta) + \lambda \|\Theta\|_1)$$

where S is the empirical covariance matrix, tr denotes the trace, $\|\cdot\|_1$ is the L_1 norm and λ is the regularization parameter which can be a scalar or a symmetric matrix that specifies different amounts of regularization for each variable. The interested reader will find a readable introduction to Gaussian graphical models, including LASSO estimation in [28, Chapter 11]. While we could use the R package `glasso` to fit the graphical model to our transformed load data, we believe that doing so would miss the special structure of the $N_Z \times N_h$ dimensional vector which has a N_Z -dimensional spatial component and a N_h -dimensional temporal component. So, taking one more step to reduce the mismatch between the

high dimensionality of the vector and the short historical data, and relying on the intuition developed for classical statistical analysis of variance with multiple factors, we think of the $N_Z \times N_h$ -by- $N_Z \times N_h$ covariance and precision matrices as tensor products (also known as Kronecker products) of $N_Z \times N_Z$ covariance and precision matrices capturing the dependencies between the spatial components (load zones), and $N_h \times N_h$ covariance and precision matrices for the temporal components (time horizons). In particular, we seek covariance matrices $\Sigma_{N_Z \times N_Z}^{spatial}$ and $\Sigma_{N_h \times N_h}^{temporal}$ such that

$$\hat{L}^d \sim N(\mathbf{0}, \Sigma) = N\left(\mathbf{0}, \Sigma_{N_Z \times N_Z}^{spatial} \otimes \Sigma_{N_h \times N_h}^{temporal}\right).$$

Searching for a precision matrix and a covariance matrix in this form dramatically reduces the number of parameters to be estimated. In the statistical literature, these special covariance matrices are called *separable*. Testing for the presence of such a special structure is sometime possible. See for example [29], [30] and the more recent paper [31], even though none of these tests have enough power in our case because of the lack of historical data. In any case, we shall work under this assumption here, and we shall use the variation on the glasso algorithm called GEMINI which was proposed in [20] to take advantage of the separable structure.

Figure 3 shows the graph of the conditional temporal dependence in the left pane, as well as the spatial conditional dependence component provided by the GEMINI algorithm. The graph of the temporal dependencies comprises two types of links. The thick links form a perfectly ordered linear chain. It points clearly at some form of Markovian structure. However, the thinner links show a residual dependence for two time horizons, which is quite intuitive. The edges, as well as their strengths in the graph of the spatial dependencies form a pattern consistent with the contiguity of the load regions as shown in pane (c) of Figure 3.

7) The next step of the analysis is the generation of Monte Carlo samples from the model. Given the covariance matrices $\Sigma_{N_Z \times N_Z}^{spatial}$ and $\Sigma_{N_h \times N_h}^{temporal}$, we can generate, for example N_s Monte Carlo samples from the $N_Z \times N_h$ -dimensional mean zero Gaussian distribution with covariance matrix $\Sigma_{N_Z \times N_Z}^{spatial} \otimes \Sigma_{N_h \times N_h}^{temporal}$. The sample matrix has N_s rows and $N_Z \times N_h$ columns. We denote by $\hat{l}_{z,h}^j$ the j -th sample for zone z and time horizon h . Recall that these Monte Carlo samples are drawn from a jointly Gaussian distribution. To obtain samples with the right marginal distributions, we need to reverse the transformations described in steps 4) and 5), namely first apply the CDF of the standard Gaussian distribution and then the quantile function (inverse CDF) of the fitted GPD $G_{z,h}$ to each Gaussian sample. Specifically, to obtain samples, denoted by $\hat{l}_{z,h}^j, j = 1, \dots, N_s$, for zone z and time horizon h which follows the marginal distribution originally fitted to the load deviations, we compute $\hat{l}_{z,h}^j = \Phi_{G_{z,h}}^{-1}\left(\Phi\left(\hat{l}_{z,h}^j\right)\right)$. At this stage, it is important to emphasize that these $N_Z \times N_h$ components depend upon each other in the way dictated by the covariance structure found by the graphical Gaussian model. The astute reader will have certainly noticed that we are actually using a $N_Z \times N_h$ -dimensional Gaussian copula




Fig. 3. Temporal component of the conditional correlation (left), spatial component of the conditional correlation (middle), geographical map of the 8 load zones in ERCOT.

to build the dependencies between the $G_{z,h}$ marginals of this $N_Z \times N_h$ -dimensional vector.

To complete the simulation cycle, we add back the forecasts to the load deviations Monte Carlo scenarios. Figure 4 shows how those scenarios fare vis-a-vis the actual and forecasts loads for May 21, 2018.




Fig. 4. 1000 Monte Carlo scenarios generated for the next 24 hours for May 21, 2018. The red (resp. blue) line gives the actual loads (resp. forecasts) for these 24 hours. The dashed lines are 10 randomly selected scenarios. The light blue bands are formed by the traces left by the Monte Carlo scenarios after we remove the smallest 1% and the largest 1% of them.

Remark 1. We relied on a form of stationarity of the load time series as we assumed that the load deviations have a single marginal distribution, irrespective of the day of the year (e.g. whether d is in the summer or the winter) or the time horizon h . Indeed, we expect that the major trends and seasonal effects are captured by the exogenously provided forecasts which are

usually constructed from a large number of meteorological variables and Bayesian averaging of many forecasting models.

B. The Impact of Fitting Heavy Tail Distributions

Here, we discuss the impact of our decision to check for the presence of heavy tails, and to possibly fit generalized Pareto distributions to the marginal laws of the deviations. In order to do so, we provide a detailed examination of the smoothness of the Monte Carlo scenarios and the coverage they provide.

Repeated overestimation by the forecasts create instances of negative deviations, and if these instances are significant because of their sizes, the empirical distribution of the deviations can exhibit a heavy lower tail which cannot be accounted for by a Gaussian model. We illustrate such an instance in Figure 5 with the example of the load zone Coast on May 21, 2018. The left pane shows that the actual loads (red curve) hover around the lower boundary of the band created by the Monte Carlo scenarios, especially during the late hours. Recall that these scenarios are obtained by adding sample scenarios of the deviations to the forecasts (blue curve). Fitting a distribution with a heavy lower tail guarantees that the scenarios do cover the actual values. The latter should be viewed as extremes, possibly rare events, but still in the realm of the model. The right pane shows the same plot produced from Monte Carlo scenarios obtained from a Gaussian model ignoring the presence of heavy tails. While one could think that having a narrower coverage band is desirable, the fact that the actual values are so far from what the model considers to be reasonable is a major flaw of this form of model.




Fig. 5. Actual loads (red) and forecasts (blue) for the 24 hours following May 21, 2018. The gray band is the trace of 10,000 Monte Carlo scenarios. The dashed lines are the sample scenarios with the minimum mean absolute percentage error (MAPE). The left pane is from our model, the right one from a Gaussian model ignoring the possible presence of heavy tails.

In a similar fashion, we can discuss the case of underestimation by forecasts. Clearly such instances can have dire consequences for the management of the electric grid, and they are feared by system operators. Figure 6 provides an illustration which makes a resounding case for the use of heavy tail distributions. We consider load in the West region on February 11, 2018. With the same convention as before, the left pane of Figure 6 shows that the actual loads for the next 24 hours (red curve), while much higher than the forecasts (blue curve) are mostly covered by the scenarios. So even on the tail end of the possibilities given by the Monte Carlo

simulations, it can still be expected to occur with significant probability and the operator of the grid should not dismiss such a pattern. In contrast, the right pane of Figure 6 shows the result when the model does not attempt to detect and fit heavy tail distributions. Clearly the curve of actual loads is far beyond the simulation band, and an operator relying on such a model will consider the 24 actual values as a possibility with zero likelihood.




Fig. 6. Same plots as in Figure 5 for February 11, 2018 to illustrate the possible consequences of an underestimation by the forecasts.

The two instances discussed above clearly illustrate the pros of checking for the presence of heavy tails in the marginal distributions of the deviations, and when detected, of using generalized Pareto distributions to process the data.

IV. MARGINAL MODEL FOR WIND POWER

Several steps of the analysis are based on the same theoretical and practical underpinnings. Therefore, we avoid most repetitions by emphasizing only the significant differences when they are relevant.

A. Modeling Steps

Let the set of wind power plants be $\mathcal{P}$. In our dataset, the number of wind power plants is $N_{\mathcal{P}} = 264$ and for each day the wind forecasts have a $N_h = 24$ hourly time horizons. For a wind power plant $p \in \mathcal{P}$, a day d and a time horizon $h \in \{0, 1, \dots, N_h - 1\}$, we define the wind deviations $W_{p,h}^d$ as: $W_{p,h}^d = \bar{W}_{p,h}^d - \tilde{W}_{p,h}^d$, where $\bar{W}_{p,h}^d$ and $\tilde{W}_{p,h}^d$ are the actual and forecast wind power productions respectively.

For each wind power plant p and for each time horizon h , the wind deviations $W_{p,h}^d, d = 1, 2, \dots, N_d$ form a daily time series of length N_d , the number of days prior to the day for which we generate scenarios. We will analyze the marginal distributions of these $N_{\mathcal{P}} \times N_h = 264 \times 24 = 6336$ stationary univariate time series. The marginal distributions of wind power deviations do not exhibit heavy tails. Indeed, the power produced by a single wind power plant is a non-negative number bounded from above by its capacity. So for a specific wind power plant, the deviation should have a distribution with support contained in an interval $[-c, +c]$ where c is the capacity of the generation asset. Therefore, for each pair of (p, h) , instead of fitting a GPD, we use the empirical cumulative distribution function (ECDF) $\Phi_{p,h}$ to apply the probability integral transform to $W_{p,h}^d, d =$

$1, 2, \dots, N_d$ and obtain a time series with uniform marginals. We then apply the quantile function Φ^{-1} of the standard Gaussian distribution to transform the time series with uniform marginals into a time series with standard Gaussian marginals. To summarize, we transform the data in the following way: $\hat{W}_{p,h}^d = \Phi^{-1} \left(\Phi_{p,z} \left(W_{p,h}^d \right) \right)$. This takes care of the steps 3) - 5) of the strategy used for the load. Fitting a Gaussian graphical model in the same manner as we did in step 6) above with the GEMINI algorithm to capture the precision and covariance matrices leads to the analog of Figure 3 which we reproduce as Figure 7. In the estimation of the spatial components of the covariance matrix using LASSO, we chose the regularization parameter λ to be a $N_{\mathcal{P}}$ -by- $N_{\mathcal{P}}$ symmetric matrix so that the ij -th element is proportional to the physical distance between the i -th and j -th wind power plants as given by their geographic coordinates.




Fig. 7. Temporal component (left), and spatial component (right) of the conditional correlation using the plain average of 51 wind ECMWF forecasts.

B. Impact of the Quality of the Forecasts.

As mentioned in Section II, we implemented simulation models for wind using two different sets of forecasts: the median of the probabilistic forecasts and a plain average of the 51 ECMWF forecasts. We show that the temporal and spatial components of the conditional correlations using the two forecasts in Figures 7 and 8. In both cases, the spatial dependency graphs are consistent with the geographical locations of the wind production assets. Figure 7 shows that using point forecasts obtained from the ECMWF forecasts, the graphical model shows the usual *Markovian-like* temporal dependence structure given by the clear chain pattern of the temporal dependency graph. However, the temporal dependency graph derived from the probabilistic forecasts doesn't indicate a clear pattern. Moreover, it depends strongly on the choice of the LASSO regularization parameters. For example, the graph shown in Figure 8 exhibits many isolated nodes. Moreover, although the temporal dependency graph contains a chain pattern for the hours from 8 to 14, the largest correlation coefficient in the graphical model is less than 20%.

C. Monte Carlo Generation Specifics

Exactly as in the case of the load simulations, we can generate Monte Carlo scenarios for the Gaussian vectors whose distributions were identified by the graphical Gaussian model, and remove the Gaussian marginals by computing the Gaussian cumulative distribution function on the entries




Fig. 8. Temporal component (left), and spatial component (right) of the conditional correlation using the medians of the probabilistic forecasts.

of these vectors. However, transitioning to the actual power productions (recall step 7) of our discussion of the load demand) requires a modicum of care. To be precise, for each wind power plant $p \in \mathcal{P}$ and time horizon h , suppose we have generated N_s Gaussian samples (scenarios) $\hat{w}_{p,h}^j, j = 1, 2, \dots, N_s$. Following the strategy for load demand, one would restore the original marginal distribution of the scenarios by computing $w_{p,h}^j = \Phi_{p,h}^{-1} \left(\Phi \left(\hat{w}_{p,h}^j \right) \right)$. Doing so would ignore the empirical evidence demonstrated in Figure 9 for two wind assets: *the marginal distribution of the deviations strongly depends upon the actual value of the forecast.*




Fig. 9. Scatterplots of the deviations (vertical axis) against the corresponding forecasts (horizontal axis).

In order to account for this phenomenon, we restore the marginal distributions of the scenarios by using the quantile function of the *conditional distribution* of the deviations given a wind power forecast. Specifically, for each wind power plant $p \in \mathcal{P}$, we first create a dataset containing all deviations irrespective of their time horizons. Then, for each $h \in \{0, 1, \dots, N_h - 1\}$ with forecast $f_{p,h}$, we denote by $\Phi_{p,f_{p,h}}$ the empirical CDF of the deviations whose corresponding forecasts are close to the given forecast value. Now, to restore the marginal distribution of the deviations, for each j , we compute $w_{p,h}^j = \Phi_{p,f_{p,h}}^{-1} \left(\Phi \left(\hat{w}_{p,h}^j \right) \right)$. Finally, we obtain the Monte Carlo scenario j of the wind power productions by adding back the forecast $f_{p,h}$ to $w_{p,h}^j$.

For illustration purposes, we show the impact at the asset-level of the quality of the forecasts by the plots of wind power scenarios for 8 wind power plant randomly selected from the 264 wind power plants. The Monte Carlo scenarios are generated from our model using the mean of the ECMWF forecasts in Figure 10, and the median of the probabilistic forecasts in Figure 11. Because the probabilistic forecasts were constructed from a bias correction procedure applied to the

ECMWF point forecasts, the forecasts seen in Figure 11 are much closer to the actuals than in the case of Figure 10. The precision with which these forecasts predict the actual values of the wind generation is remarkable. This being also true for the historical data used to fit the models, the bands covered by the scenarios in Figure 11 are significantly narrower.




Fig. 10. 1000 Monte Carlo scenarios generated on May 21, 2018 for the next 24 hours using the means of the ECMWF forecasts as point forecasts. The red (resp. blue) line gives the actual wind power (resp. point forecast) over these 24 hours, while the dashed lines are 10 randomly selected scenarios. The light blue bands are the traces left by the Monte Carlo scenarios after we remove the smallest 1% and the largest 1% of the bunch.




Fig. 11. Same as for Figure 10 except for the fact that the scenarios are now generated using the median of the probabilistic forecasts as point forecasts.

V. MARGINAL MODEL FOR SOLAR POWER

Solar power production follows a drastically different pattern than wind power production, in no small part due to the stark differences between nocturnal and diurnal productions. Thus, we adopt a different strategy to fit a model to the data.

In line with our analysis of load demand and wind power production, we work with the deviations between actuals and forecasts, i.e., for a given day d , a solar power plant $q \in \mathcal{Q}$, the set of all solar power plants, and a time horizon $h \in \{0, 1, \dots, N_h - 1\}$, we compute the deviation $S_{q,h}^d = \bar{S}_{q,h}^d - \tilde{S}_{q,h}^d$. The dataset consists $N_{\mathcal{Q}} = 226$ production assets and $N_h = 24$ -hour time horizon for the period 2017-01-01 to 2018-12-31.

In order to make sure that our model (and the subsequent Monte Carlo simulations) are non-anticipative (i.e. do not use a crystal ball) and accounts for the seasonality of solar power production (e.g. summer v.s. winter), for each given day d of 2018, we collect the historical data comprising the n days preceding d , and the days of 2017 which fall in the $2n+1$ days window centered around day d in 2017. We use this approach for any day d in 2018, and rely on the historical data identified in this way to fit the model from which we generate Monte Carlo scenarios for day d . In typical experiments we use the default value $n = 50$. This is reasonable, even when d is early in the year. For instance, if we want scenarios for 2018-07-01, the historical data points to fit the model would be 2017-05-12 to 2017-08-20 and from 2018-05-12 to 2018-06-30.

For each solar power plant q and time horizon h , to transform the data so that the marginal distributions become Gaussian, we first use the empirical CDF $\Phi_{q,h}$ to produce time series with uniform marginals, and then the standard Gaussian quantile function Φ^{-1} to end up with Gaussian marginal distributions $\hat{S}_{q,h}^d$, i.e., $\hat{S}_{q,h}^d = \Phi^{-1}(\Phi_{q,h}(S_{q,h}^d))$.

A serious hindrance to the development of joint models including solar production is its diurnal nature: the long swatches of zeros appearing during the night hours lead to singular correlation matrices which cannot be inverted for simulation purposes. The remedy we propose is to change basis for the vectors of hourly deviations and replace them by the coefficients in their decompositions on an eigenvector orthonormal basis obtained by Principal Component Analysis (PCA). Since a change of basis is a linear operation, the Gaussian nature of the data will not be affected. Moreover, as a side effect, PCA provides significant dimension reduction as we can safely ignore coefficients that do not contribute significantly to the explanation of the variation in the data.

We now explain how we practically perform the PCA of the transformed data. Denoting by N_d the number of days used to fit the model, we aggregate the transformed data over all historical days and power plants to end up with a $226 \cdot N_d \times 24$ matrix, one column for each time horizon, and one row for each pair of (power plant, day). Recall that the data comprise 226 solar power plants. We then compute the PCA orthonormal basis of 24 loadings, and a $226 \cdot N_d \times 24$ matrix of components of the original data onto the loadings basis. Given the proportions of the variation in the data explained by the successive loadings, we chose to keep, for each of the $226 \cdot N_d$ couples (power plant, day), only the first k entries in such a way that the proportion of the variance explained is greater than a percentage p_{var} . In the experiments reported in this paper we chose $p_{var} = 90\%$. The switch from the original data to their first k components is intended to resolve the issue of the diurnal zero production, more than for its *data compression* effect. We then reshape the data into $N_{\mathcal{Q}} \times k$ time series of length N_d , and standardize the data to have zero mean and unit variance. Next, we fit a graphical Gaussian model using GEMINI to obtain a $N_{\mathcal{Q}}$ -by- $N_{\mathcal{Q}}$ and a k -by- k covariance matrices that capture the dependencies among the spatial components and between the PCA loadings.

To generate scenarios, we first draw Monte Carlo samples

from the mean zero Gaussian distribution with covariance matrix defined as the tensor product of the two covariance matrices computed from the graphical model. We then reconstruct 24-dimensional vectors using the orthonormal basis of the PCA loadings. We complete the simulation by restoring the marginal distribution of the solar deviations using the strategy developed for wind (see Section IV-C) and adding back the forecasts. Figure 13 gives the graph of the spatial conditional temporal dependence. Again, it is consistent with the geographical locations of the solar production assets.

Figure 12 shows an example of asset-level solar production scenarios for 8 solar power plants randomly selected from the 226 solar power plants. As we can see, the values of the solar production given by the scenarios are only non-zero during the day. This is the result of our use of the projection of the data on the first k loadings of the PCA.




Fig. 12. 1000 Monte Carlo scenarios generated on May 21, 2018 for the next 24 hours using the plain average of the ECMWF forecasts as point forecasts. The red (resp. blue) line gives the actual solar power (resp. forecast) for these 24 hours. The dashed lines are 10 randomly selected scenarios. The light blue bands are the traces left by the Monte Carlo scenarios after we remove the smallest 1% and the largest 1% of the bunch.




Fig. 13. Spatial component of the conditional correlation for the analysis of the solar deviations using PCA.

VI. CONSOLIDATED MODEL AND SIMULATION ENGINE

Once the three marginal models discussed earlier are fitted, and the corresponding simulation engines have been devel-

oped, a joint model and a joint simulation engine can be obtained following the steps described in this section.

A. In-sample vs Out-of-sample Modeling and Simulation

To guarantee that our model (and the subsequent Monte Carlo simulations) are non-anticipative, for each given day d of 2018, we fit our model to the bunch of historical data identified earlier in the case of solar power generation. However, if we want to use the wind point forecasts given by the median of the probabilistic forecasts provided by NREL, we do not have historical data for 2017. So when the day d is early in the year, we fit the model with data for days in a window covering d , extending in the future by n days and excluding d . So, for those values of d , the estimation of the model is *in-sample*. This is unfortunate as it should not have to be if we had enough historical data.

B. Quantifying and Implementing the Dependencies

Load data being available at the zone level only, the strategy for the joint modeling of load, wind power and solar power is based on the following steps.

- 1) The first step of the analysis is again a form of *Gaussianization* of the historical data of load, wind and solar power production deviations. We do that separately, just as we did for the marginal models;
- 2) For each of the load zones and for each time horizon h , we aggregate the wind and solar Gaussianized power production historical data at the load zone level. So for each of the 8 load zones, and for each time horizon h , we sum up all the Gaussianized time series of the individual wind asset deviations to create one single wind power deviation series for this zone. Similarly, for the solar asset deviations;
- 3) Next, for each of the load zones, we aggregate the zonal load, wind and solar Gaussianized power productions across time horizons. As solar plants only generate power during the day, we only aggregate across those hours. In particular, if we denote the time horizons during sunrise and sunset as h_{sunrise} and h_{sunset} respectively, then we sum up all the Gaussianized zonal load time series to create one single load series for that zone during that day. Similarly, we aggregate wind (and solar) to create one single series for that zone during that day.
- 4) We concatenate the load, wind and solar time series from step (3) for all zones to obtain a $3 \times 8 = 24$ dimensional time series and fit a joint Gaussian graphical model at the zone level using a plain `glasso` algorithm. For this step, we do not use the `GEMINI` variation of the `glasso` algorithm because the temporal dependency has been removed as we summed up zonal load, wind and solar across time horizons in step (3).
- 5) When running the `glasso` algorithm with load, solar power and wind power, we find that the nodes of the dependency graph corresponding to wind are disconnected from the nodes corresponding to load and solar. Since we are still in the *Gaussian world* when we perform this step of the analysis, our conclusion is that, at least at the zonal

level, wind power generation is statistically independent of the electricity demand and the solar power generation. As a result, wind is modelled separately and Monte Carlo scenarios are generated for wind power from the marginal model described in Section IV

As for diurnal load and solar power, Figure 14 provides a graphical representation of the entries of the precision matrix, while Figure 15 provides a heat-map of the actual correlation matrix of the 16-dimensional vector of diurnal aggregates of load and solar power over the 8 load zones. The dependence highlighted in Figure 14 by the links between the load zones is consistent with the geographical locations of these zones, and the fact that the node corresponding to load in the "Far-West" zone is isolated is not a surprise. Links between different solar zones are also natural. Strong links between load and solar power for the same zone are also significant and expected, for example load and solar in zone South Central, or zones East and Coast, and to a lesser degree zone South. A possible reason for the negative partial correlation (pink link between "Solar_North" and "Load_West") could be the biases we noticed in the load and solar forecasts.




Fig. 14. Graphical LASSO estimate of the conditional correlation of load and solar power in the joint model described in the text.

As expected, the corresponding correlation matrix whose heat-map is given in Figure 15 is not sparse. The bottom-right fourth of the matrix confirms the strong correlations between the zone level diurnal solar power productions. These correlations are weaker among the zonal loads. The green blocks in the bottom-left quadrant of the matrix confirm the above remarks concerning the correlations between load and power production in some regions.

C. Putting Everything Together with Conditional Simulation

Given the above discussion, in order to generate Monte Carlo scenarios for the entire system on a given day d , we can first simulate samples of 24 hours wind power generation for the 264 wind power plants using the marginal model introduced and fitted in Section IV. Next, and independently of the scenarios generated for wind power, we generate jointly correlated Monte Carlo scenarios for load and solar power according to the following steps based on conditional simulation of Gaussian processes.




Fig. 15. The correlation heatmap of the load-solar joint model.

- 1) We generate jointly Gaussian scenarios for load and zonal solar power as their daily sum of the time horizons from $h_{sunrise}$ to h_{sunset} . We do just that using the joint Gaussian model estimated from the `glasso` algorithm implemented as described above.
- 2) Next for each such scenario, we generate a 8×24 (8 for each load zone and 24 for the 24 time horizons) Gaussian load Monte Carlo scenario from the marginal distribution fitted using the GEMINI algorithm as described in Section III, conditional on the fact that the sum of the time horizons in the range $[h_{sunrise}, h_{sunset}]$ are exactly equal to the values of the load scenario generated in step (1). We do that for the 8 load zones *simultaneously* to make sure that we do not wipe out the zone dependencies!
- 3) Next for each scenario generated in step (1), we generate a 226×24 (226 for each of the solar assets, and 24 for each of the time horizons) Gaussian solar Monte Carlo scenario from the marginal distribution fitted using the PCA based algorithm described in Section V, conditional on the fact that the zonal aggregates of the time horizons in the range $[h_{sunrise}, h_{sunset}]$ are exactly equal to the values of the particular solar scenario generated in step (1). We do that for the 226 solar assets *simultaneously* to make sure that we do not wipe out the dependencies between the individual asset!
- 4) Once we have the Gaussian Monte Carlo scenarios for zonal loads and individual solar assets, we can undo the Gaussian nature of the scenarios as explained in Section III and Section V and reinstate the non-necessarily Gaussian marginal distributions.

VII. EVALUATION

A. Spatial-Temporal Correlations

In this section, we examine the performance of our stochastic model (and hence of our Monte Carlo scenario generation method) by evaluating the impact of the spatial and temporal

correlations. For the sake of comparison, we place ourselves in the Gaussian stage of the modeling procedure. For each of "load", "Solar" and "Wind", we use the logarithmic scoring metric to compare the correlation structure captured by our model with several alternative correlation structures:

- 1) *Independent*: multivariate distributions with independent marginals, i.e., the correlation matrix is an identity matrix.
- 2) *Full model*: multivariate distributions with a separable spatial-temporal correlation structure estimated by the GEMINI algorithm described in the paper.
- 3) *Temporal model*: multivariate distributions with only temporal correlations. The correlations matrix is separable and the off-diagonals entries of the spatial component are set to zero.
- 4) *Spatial model*: multivariate distributions with only spatial correlations. The correlations matrix is separable and the off-diagonals entries of the temporal component are set to zero.
- 5) *Model with sample correlation*: multivariate distributions with a correlation structure given by the sample empirical correlation matrix.

In this study, we choose the negatively oriented logarithmic score as metric to assess performance. The logarithmic scoring rule is defined as

$$S = -\frac{1}{T} \sum_{t=1}^T \ln(p(x_t))$$

where $p(x)$ is the multivariate density function, x_t is an observation and T is the number of observations. In our case, p is the density function of a mean-zero multivariate Gaussian distribution with a specific correlation structure. For any given actual values of load, solar or wind power, the same procedures in the Gaussian copula approach (e.g., the probability integral transform, the inverse probability integral transform and PCA transform (for solar)) as in the model fitting stage will be applied to obtain the observation x_t to be fed to the multivariate Gaussian density.

We fit models and compute the scores for load, solar and wind (with point forecasts obtained by the median of the probabilistic forecasts and the median of the ECMWF forecasts) everyday from 2018-01-02 to 2018-12-29. All models are fit using the data prior to the day in question, except for wind with the probabilistic forecasts in which case the model is fitted using data prior and after the day in question. We report the scores for the marginal models of load, solar and wind power in Table I. The scores for models using the sample correlations are not reported for wind and solar, because the sample correlation matrices are singular due to its high dimensionality in space and the short length of historical data.

For all test cases, the independent models are obviously of the worst quality (highest scores). For load, preserving the spatial or temporal correlations or using the sample correlations indeed improve their scores, but none is as good as our full model. For solar, the "temporal" correlation is effectively the correlation between the orthogonal PCA principal components. Therefore, preserving the "temporal" correlation for

TABLE I
LOGARITHMIC SCORES FOR THE GAUSSIAN MODELS WITH DIFFERENT CORRELATION STRUCTURES.

Model Correlation	Load	Solar	Wind (Prob)	Wind (ECMWF)
Independent	272.09	2845.63	9138.31	8964.53
Temporal	111.16	2843.48	9091.16	5174.58
Spatial	261.01	1590.40	7844.20	2865.17
Full	106.09	1588.56	7864.46	746.01
Sample	166.64			

solar should not have any major impact on the quality of the model. The scores of the spatial and full model confirm this claim. For wind with probabilistic forecasts, as shown in Figure 8 and discussed in Section IV-B, the temporal correlation is weak and does not indicate a clear pattern: the score of the temporal model does not improve significantly from the independent model. However, the spatial and full model show better scores since they preserve the spatial correlation. For wind with ECMWF forecasts, the full model shows strong temporal and spatial correlations and outperforms all other models.

B. Energy and Variogram Scores

In this section, we assess the quality of our Monte Carlo scenarios using standard evaluation metrics and comparisons with methods used in the existing literature. For comparison purposes, we choose the standard vector autoregressive (VAR) model [32] as the baseline model and use the implementation in the Python package `statsmodels v0.15.0` [33] for testing. In particular,

- 1) The VAR model is fitted using the multivariate time series data obtained by taking the difference between the actuals and forecasts (namely the forecast errors or deviations).
- 2) The number of lag is selected based on Akaike Information Criterion (AIC) for up to 24 lags.
- 3) To get the most out of the fitted VAR models, we generate scenarios one time period into the future at a time. An error term is sampled using the estimated error covariance matrix, and is added to the forecasts. We repeat this process 24 times during the day to obtain 24-hour time series scenarios.

In this study, we use negatively-oriented Energy and Variogram Scores (see [18] and reference therein) as evaluation metrics. We compute them for load and wind scenarios. For load, the scenarios are generated for each day d of 2018. The models are fitted using the historical data from 2017-01-02 to the day prior to d . The load scenarios are then normalized (divided) by the zonal historical peak load. For wind, we use the one-year probabilistic forecasts data. The wind scenarios are generated for each day d from 2018-04-01 to 2018-12-29. The models are fitted using the historical data from 2018-01-02 to the day prior to d . The wind scenarios are normalized by the capacity and then aggregated to the zone level before assessment. For both load and wind, we generate 1000 scenarios everyday. Each daily scenario has length $8 \times 24 = 192$ for 8 zones and 24 hours in a day.

TABLE II
ENERGY AND VARIOGRAM SCORES FOR LOAD AND WIND.

		PGscen		VAR	
		Energy	Variogram	Energy	Variogram
Load	mean	0.334015	0.001790	0.368734	0.001872
	std. dev.	0.135419	0.001584	0.166743	0.001657
Wind	mean	0.312814	0.001807	0.616626	0.007395
	std. dev.	0.097774	0.001019	0.385249	0.008735

We report the mean and standard deviation of the Energy and Variogram scores for the model proposed in this paper (referred as PGscen) and VAR in Table II. For both load and wind, our model achieves lower scores, which indicates better performance throughout our test days.

VIII. CONCLUSION

Motivated by the need to include stochastic factors in the highly sophisticated optimization programs providing unit commitment and economic dispatch for the operation of modern electricity grids, we proposed an original joint model for electricity demand, and wind and solar power generation capturing spatial and temporal dependencies between the renewable generation assets and load zones. We developed efficient procedures to fit the model to granular data, and implemented Monte Carlo simulation algorithms producing scenarios which can be used for applications ranging from risk assessment to future generation asset siting. For this, while electricity load can remain at the zone level, wind and solar productions need to be modeled and simulated at the individual asset level. The main thrust of our model is to capture the possible and intricate dependencies and interactions between temporal and geographical components. We consider hourly data, but our model and our simulation algorithms are agnostic to the time resolution, and can be ported to higher frequencies as long as the historical data is available.

The steps of our original approach which stand out when compared to the existing literature on the subject can be summarized in the following list of items: 1) detailed statistical analysis of the marginal distributions of the deviations between the forecasts and realized values, and use of generalized Pareto distributions when heavy tails are detected; 2) *Gaussianization* of the data by processing with the fitted marginal distribution, very much in the spirit of the use of Gaussian copulas in portfolio theory; 3) use of separable covariance structures to disentangle the temporal and spatial contributions to the correlations, fitting graphical Gaussian models and producing precision matrix proxies with the *gemini* variation on the *glasso* algorithm; 4) development of a marginal statistical model for solar power production based on Principal Component Analysis and graphical LASSO; 5) construction of a Monte Carlo scenario simulator putting the various marginal models together by hierarchical conditional simulations.

We illustrate the potential of the model and its implementation on synthetic data produced by NREL for the Texas region. Our implementation, called PGscen, is part of the GitHub

<https://github.com/PrincetonUniversity/PGscen>

IX. ACKNOWLEDGMENTS

Both authors were partially supported by ARPA-E grants DE-AR0001289 and DE-AR0001390 under the PERFORM program of the US Department of Energy. We would like to thank Mike Ludkovski, Ronnie Sircar, and Glen Swindle for enlightening conversations on the content of the paper, and Michal Grzadkowski for implementing the model presented in the paper.

REFERENCES

- [1] T. Hong and S. Fan, "Probabilistic electric load forecasting: A tutorial review," *International Journal of Forecasting*, vol. 32, pp. 914 – 938, 2016.
- [2] Y. Zhang, J. Wang, and X. Wang, "Review on probabilistic forecasting of wind power generation," *Renewable & Sustainable Energy Review*, vol. 32, pp. 255 – 270, 2014.
- [3] D. van der Meer, J. Walfen, and J. Munkhammar, "Review on probabilistic forecasting of photovoltaic power production and electricity consumption," *Renewable & Sustainable Energy Review*, vol. 81, pp. 1484 – 1512, 2018.
- [4] B. M. Hodge, C. B. Martinez-Anido, Q. Wang, E. Chartan, A. Florita, and J. Kiviluoma, "The combined value of wind and solar power forecasting improvements and electricity storage," *Applied Energy*, vol. 214, p. 1–15, 2018.
- [5] A. Al-Awami and M.A.El-Sharkawi, "Statistical characterization of wind power output for a given wind power forecast," in *North American Power Symposium (NAPS)*, 2009, pp. 1–4.
- [6] M. Souto, A. Moreira, A. Veiga, A. Street, J. D. Garcia, and C. Epprecht, "A high-dimensional VARX model to simulate monthly renewable energy supply," *IEEE*, 2014, p. 1–7.
- [7] Q.Tu, S. Miao, F. Yao, Y. Li, H. Yin, J. Han, D. Zhang, and W. Yang, "Forecasting scenario generation for multiple wind farms considering time-series characteristics and spatial-temporal correlation," *Journal of Modern Power Systems and Clean Energy*, vol. 9, pp. 532–562, 2021.
- [8] C. Lucheronia, J. Bolandb, and C. Ragnoa, "Scenario generation and probabilistic forecasting analysis of spatio-temporal wind speed series with multivariate autoregressive volatility models," *Applied Energy*, vol. 239, pp. 1236 – 1241, 2019.
- [9] A. Street, D. A. Lima, A. Veiga, B. F'anzeres, L. Freire, and B. Amaral, "Fostering wind power penetration into the Brazilian forward-contract market," in *2012 IEEE Power and Energy Society General Meeting*, IEEE, 2012, p. 1–8.
- [10] J. Morales, R. Minguez, and A. Conejo, "A methodology to generate statistically dependent wind speed scenarios," *Applied Energy*, vol. 87, pp. 843–855, 2010.
- [11] D. Lee and R. Baldick, "Load and wind power scenario generation through the generalized dynamic factor model," *IEEE Transactions on Power Systems*, vol. 32, pp. 400 – 410, 2017.
- [12] R. J. Hyndman and G. Athanasopoulos, *Forecasting: Principles and Practice*, 2nd edition. OTexts, 2018.
- [13] P. Pinson, H. Madsen, H. Nielsen, G. Papaefthymiou, and B. Klöckl, "From probabilistic forecasts to statistical scenarios of short-term wind power production," *Wind Energy*, vol. 12, pp. 51–62, 2009.
- [14] J. Tastu, P. Pinson, E. Kotwa, H. Madsen, and H. Nielsen, "Spatio-temporal analysis and modeling of short-term wind power forecast errors," *Wind Energy*, vol. 14, pp. 43–60, 2011.
- [15] F. Golestaneh, H. Gooi, and P. Pinson, "Generation and evaluation of space-time trajectories of photovoltaic power," *Applied Energy*, vol. 176, pp. 80–91, 2016.
- [16] J. Tastu, P. Pinson, and H. Madsen, "Space-time trajectories of wind power generation: Parameterized precision matrices under a gaussian copula," in *Modeling and Stochastic Learning for Forecasting in High Dimensions*, ser. Lecture Notes in Statistics. Springer Verlag, 2015, vol. 217, pp. 267 – 296.
- [17] Z. B. Bouallegue, T. Heppelmann, S. Theis, and P. Pinson, "Generation of scenarios from calibrated ensemble forecasts with a dual-ensemble copula-coupling approach," *Monthly Weather Review*, vol. 144, pp. 4737 – 4750, December 2016.
- [18] A. Staid, J. Watson, R. Wets, and D. Woodruff, "Generating short-term probabilistic wind power scenarios via nonparametric forecast error density estimators," *Wind Energy*, vol. 20, pp. 837 – 848, 2017.
- [19] P. Pinson and R. Girard, "Evaluating the quality of scenarios of short-term wind power generation," *Applied Energy*, vol. 96, pp. 12–20, 2012.

- [20] S. Zhou, “GEMINI: Graph estimation with matrix variate normal instances,” *Annals of Statistics*, vol. 42, pp. 532–562, 2014.
- [21] X. Y. R. Carmona, “Glasso model for electric load and wind power and monte carlo scenario generation,” Princeton University, Tech. Rep., 2021.
- [22] ORFEUS, “Pgscen, a python package for load, wind and solar power simulation,” GitHub repository, 2021. [Online]. Available: <https://github.com/PrincetonUniversity/PGscen>
- [23] NREL, “Arpa-e perform ercot dataset produced by nrel,” GitHub repository, 2021. [Online]. Available: <https://github.com/PERFORM-Forecasts/documentation>
- [24] A. E. J. J. D. K. J. L. B. M. M. O. E. P. A. S. G. S. J. W. C. Barrows, A. Bloom, “The ieee reliability test system: A proposed 2019 update,” *IEEE Transactions on Power Systems*, vol. 35, no. 1, pp. 119 – 127, January 2020.
- [25] ERCOT, “Ercot load,” GitHub repository, 2021. [Online]. Available: <https://www.ercot.com/gridinfo/load>
- [26] R. Carmona, *Statistical Analysis of Financial Data in R, Second Edition*, ser. Springer Texts in Statistics. Springer Verlag, 2014.
- [27] —, “Rsafd,” <https://rdrr.io/rforge/Rsafd/>, 2014.
- [28] M. Wainwright, *High-Dimensional Statistics: A Non-asymptotic Viewpoint*, ser. Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge University Press, 2019.
- [29] M. Mitchell, M. Genton, and M. Gumpertz, “A likelihood ration test for separability of covariances,” *Journal of Multivariate Analysis*, vol. 97, pp. 1025 – 1043, 2006.
- [30] B. Li, M. Genton, and M. Sherman, “Testing the covariance structure of multivariate random fields,” *Biometrika*, vol. 95, pp. 813 – 829, 2008.
- [31] P. Guggenberger, F. Kleibergen, and S. Mavroeidis, “A test for Kron-necker product structure covariance matrix,” *Journal of Econometrics*, vol. available on line, 2022.
- [32] H. Lütkepohl, *New introduction to multiple time series analysis*. Springer Science & Business Media, 2005.
- [33] S. Seabold and J. Perktold, “statsmodels: Econometric and statistical modeling with python,” in *9th Python in Science Conference*, 2010.

René A. Carmona is the Paul M. Wythes ’55 Professor of Engineering and Finance in the Department of Operations Research and Financial Engineering at Princeton University. Beyond image and signal analysis, his interests include statistical data science, stochastic processes and game theory. He is a fellow of IMS, SIAM and AMS and a member of the Bachelier Society and the IEEE Power and Energy Society.

Xinshuo Yang is currently a postdoctoral research associate of Operations Research and Financial Engineering at Princeton University. Previously he was a postdoctoral fellow at Colorado School of Mines, and a postdoctoral associate at National Renewable Energy Laboratory (NREL). He received his Ph.D. in Applied Mathematics from University of Colorado Boulder.